An Introduction to libuv

Release 1.0.0

Nikhil Marathe

October 10, 2012

CONTENTS

INTRODUCTION

This 'book' is a small set of tutorials about using libuv as a high performance evented I/O library which offers the same API on Windows and Unix.

It is meant to cover the main areas of libuv, but is not a comprehensive reference discussing every function and data structure. The official libuv documentation is included directly in the libuv header file.

This book is still a work in progress, so sections may be incomplete, but I hope you will enjoy it as it grows.

1.1 Who this book is for

If you are reading this book, you are either:

1. a systems programmer, creating low-level programs such as daemons or network services and clients. You have found that the event loop approach is well suited for your application and decided to use libuv.

2. a node.js module writer, who wants to wrap platform APIs written in C or C++ with a set of (a)synchronous APIs that are exposed to JavaScript. You will use libuv purely in the context of node.js. For this you will require some other resources as the book does not cover parts specific to v8/node.js.

This book assumes that you are comfortable with the C programming language.

1.2 Background

The node.js project began in 2009 as a JavaScript environment decoupled from the browser. Using Google's V8 and Marc Lehmann's libev, node.js combined a model of I/O – evented – with a language that was well suited to the style of programming; due to the way it had been shaped by browsers. As node.js grew in popularity, it was important to make it work on Windows, but libev ran only on Unix. The Windows equivalent of kernel event notification mechanisms like kqueue or (e)poll is IOCP. libuv is an abstraction around libev or IOCP depending on the platform, providing users an API based on libev.

1.3 Code

All the code from this book is included as part of the source of the book on Github. Clone/Download the book and run `make` in the `code/` folder to compile all the examples. This book and the code is based on libuv version *node-v0.8.0* and a version is included in the `libuv/` folder which will be compiled automatically.

TWO

BASICS OF LIBUV

libuv enforces an **asynchronous**, **event-driven** style of programming. Its core job is to provide an event loop and callback based notifications of I/O and other activities. libuv offers core utilities like timers, non-blocking networking support, asynchronous file system access, child processes and more.

2.1 Event loops

In event-driven programming, an application expresses interest in certain events and respond to them when they occur. The responsibility of gathering events from the operating system or monitoring other sources of events is handled by libuv, and the user can register callbacks to be invoked when an event occurs. The event-loop usually keeps running *forever*. In pseudocode:

```
while there are still events to process:
    e = get the next event
    if there is a callback associated with e:
        call the callback
```

Some examples of events are:

- File is ready for writing
- A socket has data ready to be read
- A timer has timed out

This event loop is encapsulated by `uv_run()` – the end-all function when using libuv.

The most common activity of systems programs is to deal with input and output, rather than a lot of number-crunching. The problem with using conventional input/output functions (`read`, `fprintf`, etc.) is that they are **blocking**. The actual write to a hard disk or reading from a network, takes a disproportionately long time compared to the speed of the processor. The functions don't return until the task is done, so that your program is doing nothing. For programs which require high performance this is a major roadblock as other activities and other I/O operations are kept waiting.

One of the standard solutions is to use threads. Each blocking I/O operation is started in a separate thread (or in a thread pool). When the blocking function gets invoked in the thread, the processor can schedule another thread to run, which actually needs the CPU.

The approach followed by libuv uses another style, which is the **asynchronous, non-blocking** style. Most modern operating systems provide event notification subsystems. For example, a normal `read` call on a socket would block until the sender actually sent something. Instead, the application can request the operating system to watch the socket and put an event notification in the queue. The application can inspect the events at its convenience (perhaps doing some number crunching before to use the processor to the maximum) and grab the data. It is **asynchronous** because the application expressed interest at one point, then used the data at another point (in time and space). It is **non-blocking** because the application process was free to do other tasks. This fits in well with libuv's event-loop approach,

since the operating system events can be treated as just another libuv event. The non-blocking ensures that other events can continue to be handled as fast they come in [1].

Note: How the I/O is run in the background is not of our concern, but due to the way our computer hardware works, with the thread as the basic unit of the processor, libuv and OSes will usually run background/worker threads and/or polling to perform tasks in a non-blocking manner.

Bert Belder, one of the libuv core developers has a small video explaining the architecture of libuv and its background. If you have no prior experience with either libuv or libev, it is a quick, useful watch.

2.2 Hello World

With the basics out of the way, lets write our first libuv program. It does nothing, except start a loop which will exit immediately.

helloworld/main.c

```c
#include <stdio.h>
#include <uv.h>

int main() {
    uv_loop_t *loop = uv_loop_new();

    printf("Now quitting.\n");
    uv_run(loop);

    return 0;
}
```

This program quits immediately because it has no events to process. A libuv event loop has to be told to watch out for events using the various API functions.

2.2.1 Default loop

A default loop is provided by libuv and can be accessed using `uv_default_loop()`. You should use this loop if you only want a single loop.

Note: node.js uses the default loop as its main loop. If you are writing bindings you should be aware of this.

2.3 Watchers

Watchers are how users of libuv express interest in particular events. Watchers are opaque structs named as `uv_TYPE_t` where type signifies what the watcher is used for. A full list of watchers supported by libuv is:

---

[1] Depending on the capacity of the hardware of course.

libuv watchers

```
typedef struct uv_loop_s uv_loop_t;
typedef struct uv_ares_task_s uv_ares_task_t;
typedef struct uv_err_s uv_err_t;
typedef struct uv_handle_s uv_handle_t;
typedef struct uv_stream_s uv_stream_t;
typedef struct uv_tcp_s uv_tcp_t;
typedef struct uv_udp_s uv_udp_t;
typedef struct uv_pipe_s uv_pipe_t;
typedef struct uv_tty_s uv_tty_t;
typedef struct uv_poll_s uv_poll_t;
typedef struct uv_timer_s uv_timer_t;
typedef struct uv_prepare_s uv_prepare_t;
typedef struct uv_check_s uv_check_t;
typedef struct uv_idle_s uv_idle_t;
typedef struct uv_async_s uv_async_t;
typedef struct uv_getaddrinfo_s uv_getaddrinfo_t;
typedef struct uv_process_s uv_process_t;
typedef struct uv_counters_s uv_counters_t;
typedef struct uv_cpu_info_s uv_cpu_info_t;
```

Note: All watcher structs are subclasses of `uv_handle_t` and often referred to as **handles** in libuv and in this text.

Watchers are setup by a corresponding:

```
uv_TYPE_init(uv_TYPE_t*)
```

function.

Note: Some watcher initialization functions require the loop as a first argument.

A watcher is set to actually listen for events by invoking:

```
uv_TYPE_start(uv_TYPE_t*, callback)
```

and stopped by calling the corresponding:

```
uv_TYPE_stop(uv_TYPE_t*)
```

Callbacks are functions which are called by libuv whenever an event the watcher is interested in has taken place. Application specific logic will usually be implemented in the callback. For example, an IO watcher's callback will receive the data read from a file, a timer callback will be triggered on timeout and so on.

2.3.1 Idling

Here is an example of using a watcher. An idle watcher's callback is repeatedly called. There are some deeper semantics, discussed in *Utilities*, but we'll ignore them for now. Let's just use an idle watcher to look at the watcher life cycle and see how `uv_run()` will now block because a watcher is present. The idle watcher is stopped when the count is reached and `uv_run()` exits since no event watchers are active.

idle-basic/main.c

```c
#include <stdio.h>
#include <uv.h>

int64_t counter = 0;

void wait_for_a_while(uv_idle_t* handle, int status) {
    counter++;

    if (counter >= 10e6)
        uv_idle_stop(handle);
}

int main() {
    uv_idle_t idler;

    uv_idle_init(uv_default_loop(), &idler);
    uv_idle_start(&idler, wait_for_a_while);

    printf("Idling...\n");
    uv_run(uv_default_loop());

    return 0;
}
```

void *data pattern

note about not necessarily creating type structs on the stack

---

FILESYSTEM

Simple filesystem read/write is achieved using the `uv_fs_*` functions and the `uv_fs_t` struct.

Note: The libuv filesystem operations are different from *socket operations*. Socket operations use the non-blocking operations provided by the operating system. Filesystem operations use blocking functions internally, but invoke these functions in a thread pool and notify watchers registered with the event loop when application interaction is required.

Note: The fs operations are actually a part of libeio on Unix systems. `libeio` is a separate library written by the author of `libev`.

All filesystem functions have two forms - *synchronous* and *asynchronous*.

The *synchronous* forms automatically get called (and **block**) if no callback is specified. The return value of functions is the equivalent Unix return value (usually 0 on success, -1 on error).

The *asynchronous* form is called when a callback is passed and the return value is 0.

3.1 Reading/Writing files

A file descriptor is obtained using

```
int uv_fs_open(uv_loop_t* loop, uv_fs_t* req, const char* path, int flags, int mode, uv_fs_cb cb)
```

`flags` and `mode` are standard Unix flags. libuv takes care of converting to the appropriate Windows flags.

File descriptors are closed using

```
int uv_fs_close(uv_loop_t* loop, uv_fs_t* req, uv_file file, uv_fs_cb cb)
```

Filesystem operation callbacks have the signature:

```
void callback(uv_fs_t* req);
```

Let's see a simple implementation of `cat`. We start with registering a callback for when the file is opened:

uvcat/main.c - opening a file

```
1   void on_open(uv_fs_t *req) {
2       if (req->result != -1) {
3           uv_fs_read(uv_default_loop(), &read_req, req->result,
4                      buffer, sizeof(buffer), -1, on_read);
5       }
6       else {
7           fprintf(stderr, "error opening file: %d\n", req->errorno);
8       }
9       uv_fs_req_cleanup(req);
10  }
```

The `result` field of a `uv_fs_t` is the file descriptor in case of the `uv_fs_open` callback. If the file is successfully opened, we start reading it.

> **Warning:** The `uv_fs_req_cleanup()` function must be called to free internal memory allocations in libuv.

uvcat/main.c - read callback

```
1   void on_read(uv_fs_t *req) {
2       uv_fs_req_cleanup(req);
3       if (req->result < 0) {
4           fprintf(stderr, "Read error: %s\n", uv_strerror(uv_last_error(uv_default_loop())));
5       }
6       else if (req->result == 0) {
7           uv_fs_t close_req;
8           // synchronous
9           uv_fs_close(uv_default_loop(), &close_req, open_req.result, NULL);
10      }
11      else {
12          uv_fs_write(uv_default_loop(), &write_req, 1, buffer, req->result, -1, on_write);
13      }
14  }
```

In the case of a read call, you should pass an *initialized* buffer which will be filled with data before the read callback is triggered.

In the read callback the `result` field is 0 for EOF, -1 for error and the number of bytes read on success.

Here you see a common pattern when writing asynchronous programs. The `uv_fs_close()` call is performed synchronously. *Usually tasks which are one-off, or are done as part of the startup or shutdown stage are performed synchronously, since we are interested in fast I/O when the program is going about its primary task and dealing with multiple I/O sources*. For solo tasks the performance difference usually is negligible and may lead to simpler code.

We can generalize the pattern that the actual return value of the original system call is stored in `uv_fs_t.result`.

Filesystem writing is similarly simple using `uv_fs_write()`. *Your callback will be triggered after the write is complete*. In our case the callback simply drives the next read. Thus read and write proceed in lockstep via callbacks.

uvcat/main.c - write callback

```
1   void on_write(uv_fs_t *req) {
2       uv_fs_req_cleanup(req);
3       if (req->result < 0) {
4           fprintf(stderr, "Write error: %s\n", uv_strerror(uv_last_error(uv_default_loop())));
5       }
```

```
6       else {
7           uv_fs_read(uv_default_loop(), &read_req, open_req.result, buffer, sizeof(buffer), -1, on_read
8       }
9   }
```

Note: The error usually stored in errno can be accessed from uv_fs_t.errorno, but converted to a standard UV_* error code. There is currently no way to directly extract a string error message from the errorno field.

Warning: Due to the way filesystems and disk drives are configured for performance, a write that 'succeeds' may not be committed to disk yet. See uv_fs_fsync for stronger guarantees.

We set the dominos rolling in main():

uvcat/main.c

```
1   int main(int argc, char **argv) {
2       uv_fs_open(uv_default_loop(), &open_req, argv[1], O_RDONLY, 0, on_open);
3       uv_run(uv_default_loop());
4       return 0;
5   }
```

3.2 Filesystem operations

All the standard filesystem operations like unlink, rmdir, stat are supported asynchronously and have intuitive argument order. They follow the same patterns as the read/write/open calls, returning the result in the uv_fs_t.result field. The full list:

Filesystem operations

```
UV_EXTERN int uv_fs_unlink(uv_loop_t* loop, uv_fs_t* req, const char* path,
    uv_fs_cb cb);

UV_EXTERN int uv_fs_write(uv_loop_t* loop, uv_fs_t* req, uv_file file,
    void* buf, size_t length, int64_t offset, uv_fs_cb cb);

UV_EXTERN int uv_fs_mkdir(uv_loop_t* loop, uv_fs_t* req, const char* path,
    int mode, uv_fs_cb cb);

UV_EXTERN int uv_fs_rmdir(uv_loop_t* loop, uv_fs_t* req, const char* path,
    uv_fs_cb cb);

UV_EXTERN int uv_fs_readdir(uv_loop_t* loop, uv_fs_t* req,
    const char* path, int flags, uv_fs_cb cb);

UV_EXTERN int uv_fs_stat(uv_loop_t* loop, uv_fs_t* req, const char* path,
    uv_fs_cb cb);

UV_EXTERN int uv_fs_fstat(uv_loop_t* loop, uv_fs_t* req, uv_file file,
    uv_fs_cb cb);
```

```
UV_EXTERN int uv_fs_rename(uv_loop_t* loop, uv_fs_t* req, const char* path,
    const char* new_path, uv_fs_cb cb);

UV_EXTERN int uv_fs_fsync(uv_loop_t* loop, uv_fs_t* req, uv_file file,
    uv_fs_cb cb);

UV_EXTERN int uv_fs_fdatasync(uv_loop_t* loop, uv_fs_t* req, uv_file file,
    uv_fs_cb cb);

UV_EXTERN int uv_fs_ftruncate(uv_loop_t* loop, uv_fs_t* req, uv_file file,
    int64_t offset, uv_fs_cb cb);

UV_EXTERN int uv_fs_sendfile(uv_loop_t* loop, uv_fs_t* req, uv_file out_fd,
    uv_file in_fd, int64_t in_offset, size_t length, uv_fs_cb cb);

UV_EXTERN int uv_fs_chmod(uv_loop_t* loop, uv_fs_t* req, const char* path,
    int mode, uv_fs_cb cb);

UV_EXTERN int uv_fs_utime(uv_loop_t* loop, uv_fs_t* req, const char* path,
    double atime, double mtime, uv_fs_cb cb);

UV_EXTERN int uv_fs_futime(uv_loop_t* loop, uv_fs_t* req, uv_file file,
    double atime, double mtime, uv_fs_cb cb);

UV_EXTERN int uv_fs_lstat(uv_loop_t* loop, uv_fs_t* req, const char* path,
    uv_fs_cb cb);

UV_EXTERN int uv_fs_link(uv_loop_t* loop, uv_fs_t* req, const char* path,
    const char* new_path, uv_fs_cb cb);

/*
 * This flag can be used with uv_fs_symlink on Windows
 * to specify whether path argument points to a directory.
 */
#define UV_FS_SYMLINK_DIR          0x0001

/*
 * This flag can be used with uv_fs_symlink on Windows
 * to specify whether the symlink is to be created using junction points.
 */
#define UV_FS_SYMLINK_JUNCTION     0x0002

UV_EXTERN int uv_fs_symlink(uv_loop_t* loop, uv_fs_t* req, const char* path,
    const char* new_path, int flags, uv_fs_cb cb);

UV_EXTERN int uv_fs_readlink(uv_loop_t* loop, uv_fs_t* req, const char* path,
    uv_fs_cb cb);

UV_EXTERN int uv_fs_fchmod(uv_loop_t* loop, uv_fs_t* req, uv_file file,
    int mode, uv_fs_cb cb);

UV_EXTERN int uv_fs_chown(uv_loop_t* loop, uv_fs_t* req, const char* path,
    int uid, int gid, uv_fs_cb cb);

UV_EXTERN int uv_fs_fchown(uv_loop_t* loop, uv_fs_t* req, uv_file file,
    int uid, int gid, uv_fs_cb cb);
```

3.3 Buffers and Streams

The basic I/O tool in libuv is the stream (`uv_stream_t`). TCP sockets, UDP sockets, and pipes for file I/O and IPC are all treated as stream subclasses.

Streams are initialized using custom functions for each subclass, then operated upon using

```
int uv_read_start(uv_stream_t*, uv_alloc_cb alloc_cb, uv_read_cb read_cb);
int uv_read_stop(uv_stream_t*);
int uv_write(uv_write_t* req, uv_stream_t* handle,
             uv_buf_t bufs[], int bufcnt, uv_write_cb cb);
```

The stream based functions are simpler to use than the filesystem ones and libuv will automatically keep reading from a stream when `uv_read_start()` is called once, until `uv_read_stop()` is called.

The discrete unit of data is the buffer – `uv_buf_t`. This is simply a collection of a pointer to bytes (`uv_buf_t.base`) and the length (`uv_buf_t.len`). The `uv_buf_t` is lightweight and passed around by value. What does require management is the actual bytes, which have to be allocated and freed by the application.

To demonstrate streams we will need to use `uv_pipe_t`. This allows streaming local files [1]. Here is a simple tee utility using libuv. Doing all operations asynchronously shows the power of evented I/O. The two writes won't block each other, but we've to be careful to copy over the buffer data to ensure we don't free a buffer until it has been written.

The program is to be executed as:

```
./uvtee <output_file>
```

We start of opening pipes on the files we require. libuv pipes to a file are opened as bidirectional by default.

uvtee/main.c - read on pipes

```
1  int main(int argc, char **argv) {
2      loop = uv_default_loop();
3
4      uv_pipe_init(loop, &stdin_pipe, 0);
5      uv_pipe_open(&stdin_pipe, 0);
6
7      uv_pipe_init(loop, &stdout_pipe, 0);
8      uv_pipe_open(&stdout_pipe, 1);
9
10     uv_fs_t file_req;
11     int fd = uv_fs_open(loop, &file_req, argv[1], O_CREAT | O_RDWR, 0644, NULL);
12     uv_pipe_init(loop, &file_pipe, 0);
13     uv_pipe_open(&file_pipe, fd);
14
15     uv_read_start((uv_stream_t*)&stdin_pipe, alloc_buffer, read_stdin);
16
17     uv_run(loop);
18     return 0;
19  }
```

The third argument of `uv_pipe_init()` should be set to 1 for IPC using named pipes. This is covered in *Processes*. The `uv_pipe_open()` call associates the file descriptor with the file.

We start monitoring `stdin`. The `alloc_buffer` callback is invoked as new buffers are required to hold incoming data. `read_stdin` will be called with these buffers.

---

[1] see *Pipes*

uvtee/main.c - reading buffers

```
1  uv_buf_t alloc_buffer(uv_handle_t *handle, size_t suggested_size) {
2      return uv_buf_init((char*) malloc(suggested_size), suggested_size);
3  }
4
5  void read_stdin(uv_stream_t *stream, ssize_t nread, uv_buf_t buf) {
6      if (nread == -1) {
7          if (uv_last_error(loop).code == UV_EOF) {
8              uv_close((uv_handle_t*)&stdin_pipe, NULL);
9              uv_close((uv_handle_t*)&stdout_pipe, NULL);
10             uv_close((uv_handle_t*)&file_pipe, NULL);
11         }
12     }
13     else {
14         if (nread > 0) {
15             write_data((uv_stream_t*)&stdout_pipe, nread, buf, on_stdout_write);
16             write_data((uv_stream_t*)&file_pipe, nread, buf, on_file_write);
17         }
18     }
19     if (buf.base)
20         free(buf.base);
21 }
```

The standard `malloc` is sufficient here, but you can use any memory allocation scheme. For example, node.js uses its own slab allocator which associates buffers with V8 objects.

The read callback `nread` parameter is -1 on any error. This error might be EOF, in which case we close all the streams, using the generic close function `uv_close()` which deals with the handle based on its internal type. Otherwise `nread` is a non-negative number and we can attempt to write that many bytes to the output streams. Finally remember that buffer allocation and deallocation is application responsibility, so we free the data.

uvtee/main.c - Write to pipe

```
1  typedef struct {
2      uv_write_t req;
3      uv_buf_t buf;
4  } write_req_t;
5
6  void free_write_req(uv_write_t *req) {
7      write_req_t *wr = (write_req_t*) req;
8      free(wr->buf.base);
9      free(wr);
10 }
11
12 void on_stdout_write(uv_write_t *req, int status) {
13     free_write_req(req);
14 }
15
16 void on_file_write(uv_write_t *req, int status) {
17     free_write_req(req);
18 }
19
20 void write_data(uv_stream_t *dest, size_t size, uv_buf_t buf, uv_write_cb callback) {
21     write_req_t *req = (write_req_t*) malloc(sizeof(write_req_t));
22     req->buf = uv_buf_init((char*) malloc(size), size);
23     memcpy(req->buf.base, buf.base, size);
```

```
24        uv_write((uv_write_t*) req, (uv_stream_t*)dest, &req->buf, 1, callback);
25    }
```

write_data() makes a copy of the buffer obtained from read. Again, this buffer does not get passed through to the callback trigged on write completion. To get around this we wrap a write request and a buffer in write_req_t and unwrap it in the callbacks.

> **Warning:** If your program is meant to be used with other programs it may knowingly or unknowingly be writing to a pipe. This makes it susceptible to aborting on receiving a SIGPIPE. It is a good idea to insert:
>
> signal(SIGPIPE, SIG_IGN)
>
> in the initialization stages of your application.

3.4 File change events

All modern operating systems provide APIs to put watches on individual files or directories and be informed when the files are modified. libuv wraps common file change notification libraries [2]. This is one of the more inconsistent parts of libuv. File change notification systems are themselves extremely varied across platforms so getting everything working everywhere is difficult. To demonstrate, I'm going to build a simple utility which runs a command whenever any of the watched files change:

```
./onchange <command> <file1> [file2] ...
```

The file change notification is started using uv_fs_event_init():

onchange/main.c - The setup

```
1        while (argc-- > 2) {
2            fprintf(stderr, "Adding watch on %s\n", argv[argc]);
3            uv_fs_event_init(loop, (uv_fs_event_t*) malloc(sizeof(uv_fs_event_t)), argv[argc], run_comman
4        }
```

The third argument is the actual file or directory to monitor. The last argument, flags, can be:

```
enum uv_fs_event {
  UV_RENAME = 1,
  UV_CHANGE = 2
};
```

but both are currently unimplemented on all platforms.

> **Warning:** You will in fact raise an assertion error if you pass any flags. So stick to 0.

The callback will receive the following arguments:

1. uv_fs_event_t *handle - The watcher. The filename field of the watcher is the file on which the watch was set.

2. const char *filename - If a directory is being monitored, this is the file which was changed. Only non-null on Linux and Windows. May be null even on those platforms.

---

[2] inotify on Linux, kqueue on BSDs, ReadDirectoryChangesW on Windows, event ports on Solaris, unsupported on Cygwin

3. `int flags` - one of `UV_RENAME` or `UV_CHANGE`.

4. `int status` - Currently 0.

In our example we simply print the arguments and run the command using `system()`.

onchange/main.c - file change notification callback

```c
void run_command(uv_fs_event_t *handle, const char *filename, int events, int status) {
    fprintf(stderr, "Change detected in %s: ", handle->filename);
    if (events == UV_RENAME)
        fprintf(stderr, "renamed");
    if (events == UV_CHANGE)
        fprintf(stderr, "changed");

    fprintf(stderr, " %s\n", filename ? filename : "");
    system(command);
}
```

NETWORKING

Networking in libuv is not much different from directly using the BSD socket interface, some things are easier, all are non-blocking, but the concepts stay the same. In addition libuv offers utility functions to abstract the annoying, repetitive and low-level tasks like setting up sockets using the BSD socket structures, DNS lookup, and tweaking various socket parameters.

The `uv_tcp_t` and `uv_udp_t` structures are used for network I/O.

4.1 TCP

TCP is a connection oriented, stream protocol and is therefore based on the libuv streams infrastructure.

4.1.1 Server

Server sockets proceed by:

1. `uv_tcp_init` the TCP watcher.

2. `uv_tcp_bind` it.

3. Call `uv_listen` on the watcher to have a callback invoked whenever a new connection is established by a client.

4. Use `uv_accept` to accept the connection.

5. Use *stream operations* to communicate with the client.

Here is a simple echo server

tcp-echo-server/main.c - The listen socket

```
1  int main() {
2      loop = uv_default_loop();
3
4      uv_tcp_t server;
5      uv_tcp_init(loop, &server);
6
7      struct sockaddr_in bind_addr = uv_ip4_addr("0.0.0.0", 7000);
8      uv_tcp_bind(&server, bind_addr);
9      int r = uv_listen((uv_stream_t*) &server, 128, on_new_connection);
10     if (r) {
11         fprintf(stderr, "Listen error %s\n", uv_err_name(uv_last_error(loop)));
```

```
12          return 1;
13      }
14      return uv_run(loop);
15  }
```

You can see the utility function `uv_ip4_addr` being used to convert from a human readable IP address, port pair to the sockaddr_in structure required by the BSD socket APIs. The reverse can be obtained using `uv_ip4_name`.

Note: In case it wasn't obvious there are `uv_ip6_*` analogues for the ip4 functions.

Most of the setup functions are normal functions since its all CPU-bound. `uv_listen` is where we return to libuv's callback style. The second arguments is the backlog queue – the maximum length of queued connections.

When a connection is initiated by clients, the callback is required to set up a watcher for the client socket and associate the watcher using `uv_accept`. In this case we also establish interest in reading from this stream.

tcp-echo-server/main.c - Accepting the client

```
1  void on_new_connection(uv_stream_t *server, int status) {
2      if (status == -1) {
3          // error!
4          return;
5      }
6
7      uv_tcp_t *client = (uv_tcp_t*) malloc(sizeof(uv_tcp_t));
8      uv_tcp_init(loop, client);
9      if (uv_accept(server, (uv_stream_t*) client) == 0) {
10         uv_read_start((uv_stream_t*) client, alloc_buffer, echo_read);
11     }
12     else {
13         uv_close((uv_handle_t*) client, NULL);
14     }
15 }
```

The remaining set of functions is very similar to the streams example and can be found in the code. Just remember to call `uv_close` when the socket isn't required. This can be done even in the `uv_listen` callback if you are not interested in accepting the connection.

4.1.2 Client

Where you do bind/listen/accept, on the client side its simply a matter of calling `uv_tcp_connect`. The same `uv_connect_cb` style callback of `uv_listen` is used by `uv_tcp_connect`. Try:

```
uv_tcp_t socket;
uv_tcp_init(loop, &socket);

uv_connect_t connect;

struct sockaddr_in dest = uv_ip4_addr("127.0.0.1", 80);

uv_tcp_connect(&connect, &socket, dest, on_connect);
```

where `on_connect` will be called after the connection is established.

4.2 UDP

The User Datagram Protocol offers connectionless, unreliable network communication. Hence libuv doesn't offer a stream. Instead libuv provides non-blocking UDP support via the *uv_udp_t* (for receiving) and *uv_udp_send_t* (for sending) structures and related functions. That said, the actual API for reading/writing is very similar to normal stream reads. To look at how UDP can be used, the example shows the first stage of obtaining an IP address from a DHCP server – DHCP Discover.

Note: You will have to run *udp-dhcp* as **root** since it uses well known port numbers below 1024.

---

udp-dhcp/main.c - Setup and send UDP packets

```
1  uv_loop_t *loop;
2  uv_udp_t send_socket;
3  uv_udp_t recv_socket;
4
5  int main() {
6      loop = uv_default_loop();
7
8      uv_udp_init(loop, &recv_socket);
9      struct sockaddr_in recv_addr = uv_ip4_addr("0.0.0.0", 68);
10     uv_udp_bind(&recv_socket, recv_addr, 0);
11     uv_udp_recv_start(&recv_socket, alloc_buffer, on_read);
12
13     uv_udp_init(loop, &send_socket);
14     uv_udp_bind(&send_socket, uv_ip4_addr("0.0.0.0", 0), 0);
15     uv_udp_set_broadcast(&send_socket, 1);
16
17     uv_udp_send_t send_req;
18     uv_buf_t discover_msg = make_discover_msg(&send_req);
19
20     struct sockaddr_in send_addr = uv_ip4_addr("255.255.255.255", 67);
21     uv_udp_send(&send_req, &send_socket, &discover_msg, 1, send_addr, on_send);
22
23     return uv_run(loop);
24 }
```

Note: The IP address `0.0.0.0` is used to bind to all interfaces. The IP address `255.255.255.255` is a broadcast address meaning that packets will be sent to all interfaces on the subnet. port `0` means that the OS randomly assigns a port.

First we setup the receiving socket to bind on all interfaces on port 68 (DHCP client) and start a read watcher on it. Then we setup a similar send socket and use `uv_udp_send` to send a *broadcast message* on port 67 (DHCP server).

It is **necessary** to set the broadcast flag, otherwise you will get an `EACCES` error [1]. The exact message being sent is irrelevant to this book and you can study the code if you are interested. As usual the read and write callbacks will receive a status code of -1 if something went wrong.

Since UDP sockets are not connected to a particular peer, the read callback receives an extra parameter about the sender of the packet. The `flags` parameter may be `UV_UDP_PARTIAL` if the buffer provided by your allocator was not large enough to hold the data. *In this case the OS will discard the data that could not fit* (That's UDP for you!).

---

[1] http://beej.us/guide/bgnet/output/html/multipage/advanced.html#broadcast

udp-dhcp/main.c - Reading packets

```
1  void on_read(uv_udp_t *req, ssize_t nread, uv_buf_t buf, struct sockaddr *addr, unsigned flags) {
2      if (nread == -1) {
3          fprintf(stderr, "Read error %s\n", uv_err_name(uv_last_error(loop)));
4          uv_close((uv_handle_t*) req, NULL);
5          free(buf.base);
6          return;
7      }
8
9      char sender[17] = { 0 };
10     uv_ip4_name((struct sockaddr_in*) addr, sender, 16);
11     fprintf(stderr, "Recv from %s\n", sender);
12
13     // ... DHCP specific code
14
15     free(buf.base);
16     uv_udp_recv_stop(req);
17 }
```

4.2.1 UDP Options

Time-to-live

The TTL of packets sent on the socket can be changed using `uv_udp_set_ttl`.

IPv6 stack only

IPv6 sockets can be used for both IPv4 and IPv6 communication. If you want to restrict the socket to IPv6 only, pass the `UV_UDP_IPV6ONLY` flag to `uv_udp_bind6` [2].

Multicast

A socket can (un)subscribe to a multicast group using:

```
UV_EXTERN int uv_udp_set_membership(uv_udp_t* handle,
    const char* multicast_addr, const char* interface_addr,
    uv_membership membership);
```

where `membership` is `UV_JOIN_GROUP` or `UV_LEAVE_GROUP`.

Local loopback of multicast packets is enabled by default [3], use `uv_udp_set_multicast_loop` to switch it off.

The packet time-to-live for multicast packets can be changed using `uv_udp_set_multicast_ttl`.

4.3 Querying DNS

libuv provides asynchronous DNS resolution. For this it provides its own `getaddrinfo` replacement [4]. In the callback you can perform normal socket operations on the retrieved addresses. Let's connect to Freenode to see an

---

[2] on Windows only supported on Windows Vista and later.
[3] http://www.tldp.org/HOWTO/Multicast-HOWTO-6.html#ss6.1
[4] libuv use the system `getaddrinfo` in the libuv threadpool. libuv v0.8.0 and earlier also included c-ares as an alternative, but this has been removed in v0.9.0.

example of DNS resolution.

dns/main.c

```
1   int main() {
2       loop = uv_default_loop();
3
4       struct addrinfo hints;
5       hints.ai_family = PF_INET;
6       hints.ai_socktype = SOCK_STREAM;
7       hints.ai_protocol = IPPROTO_TCP;
8       hints.ai_flags = 0;
9
10      uv_getaddrinfo_t resolver;
11      fprintf(stderr, "irc.freenode.net is... ");
12      int r = uv_getaddrinfo(loop, &resolver, on_resolved, "irc.freenode.net", "6667", &hints);
13
14      if (r) {
15          fprintf(stderr, "getaddrinfo call error %s\n", uv_err_name(uv_last_error(loop)));
16          return 1;
17      }
18      return uv_run(loop);
19  }
```

If `uv_getaddrinfo` returns non-zero, something went wrong in the setup and your callback won't be invoked at all. All arguments can be freed immediately after `uv_getaddrinfo` returns. The *hostname*, *servname* and *hints* structures are documented in the getaddrinfo man page.

In the resolver callback, you can pick any IP from the linked list of `struct addrinfo(s)`. This also demonstrates `uv_tcp_connect`. It is necessary to call `uv_freeaddrinfo` in the callback.

dns/main.c

```
1   void on_resolved(uv_getaddrinfo_t *resolver, int status, struct addrinfo *res) {
2       if (status == -1) {
3           fprintf(stderr, "getaddrinfo callback error %s\n", uv_err_name(uv_last_error(loop)));
4           return;
5       }
6
7       char addr[17] = {'\0'};
8       uv_ip4_name((struct sockaddr_in*) res->ai_addr, addr, 16);
9       fprintf(stderr, "%s\n", addr);
10
11      uv_connect_t *connect_req = (uv_connect_t*) malloc(sizeof(uv_connect_t));
12      uv_tcp_t *socket = (uv_tcp_t*) malloc(sizeof(uv_tcp_t));
13      uv_tcp_init(loop, socket);
14
15      connect_req->data = (void*) socket;
16      uv_tcp_connect(connect_req, socket, *(struct sockaddr_in*) res->ai_addr, on_connect);
17
18      uv_freeaddrinfo(res);
19  }
```

4.4 Network interfaces

Information about the system's network interfaces can be obtained through libuv using `uv_interface_addresses`. This simple program just prints out all the interface details so you get an idea of the fields that are available. This is useful to allow your service to bind to IP addresses when it starts.

interfaces/main.c

```c
1  #include <stdio.h>
2  #include <uv.h>
3
4  int main() {
5      char buf[512];
6      uv_interface_address_t *info;
7      int count, i;
8
9      uv_interface_addresses(&info, &count);
10     i = count;
11
12     printf("Number of interfaces: %d\n", count);
13     while (i--) {
14         uv_interface_address_t interface = info[i];
15
16         printf("Name: %s\n", interface.name);
17         printf("Internal? %s\n", interface.is_internal ? "Yes" : "No");
18
19         if (interface.address.address4.sin_family == AF_INET) {
20             uv_ip4_name(&interface.address.address4, buf, sizeof(buf));
21             printf("IPv4 address: %s\n", buf);
22         }
23         else if (interface.address.address4.sin_family == AF_INET6) {
24             uv_ip6_name(&interface.address.address6, buf, sizeof(buf));
25             printf("IPv6 address: %s\n", buf);
26         }
27
28         printf("\n");
29     }
30
31     uv_free_interface_addresses(info, count);
32     return 0;
33 }
```

`is_internal` is true for loopback interfaces. Note that if a physical interface has multiple IPv4/IPv6 addresses, the name will be reported multiple times, with each address being reported once.

THREADS

Wait a minute? Why are we on threads? Aren't event loops supposed to be **the way** to do *web-scale programming*? Well no. Threads are still the medium in which the processor does its job, and threads are mighty useful sometimes, even though you might have to wade through synchronization primitives.

Threads are used internally to fake the asynchronous nature of all the system calls. libuv also uses threads to allow you, the application, to perform a task asynchronously that is actually blocking, by spawning a thread and collecting the result when it is done.

Today there are two predominant thread libraries. The Windows threads implementation and pthreads. libuv's thread API is analogous to the pthread API and often has similar semantics.

A notable aspect of libuv's thread facilities is that it is a self contained section within libuv. Whereas other features intimately depend on the event loop and callback principles, threads are complete agnostic, they block as required, signal errors directly via return values and, as shown in the *first example*, don't even require a running event loop.

libuv's thread API is also very limited since the semantics and syntax of threads are different on all platforms, with different levels of completeness.

This chapter makes the following assumption: **There is only one event loop, running in one thread (the main thread)**. No other thread interacts with the event loop (except using uv_async_send). *Multiple event loops* covers running event loops in different threads and managing them.

5.1 Core thread operations

There isn't much here, you just start a thread using uv_thread_create() and wait for it to close using uv_thread_join().

thread-create/main.c

```
1   int main() {
2       int tracklen = 10;
3       uv_thread_t hare_id;
4       uv_thread_t tortoise_id;
5       uv_thread_create(&hare_id, hare, &tracklen);
6       uv_thread_create(&tortoise_id, tortoise, &tracklen);
7
8       uv_thread_join(&hare_id);
9       uv_thread_join(&tortoise_id);
10      return 0;
11  }
```

Tip: `uv_thread_t` is just an alias for `pthread_t` on Unix, but this is an implementation detail, avoid depending on it to always be true.

The second parameter is the function which will serve as the entry point for the thread, the last parameter is a `void *` argument which can be used to pass custom parameters to the thread. The function `hare` will now run in a separate thread, scheduled pre-emptively by the operating system:

thread-create/main.c

```
1  void hare(void *arg) {
2      int tracklen = *((int *) arg);
3      while (tracklen) {
4          tracklen--;
5          sleep(1);
6          fprintf(stderr, "Hare ran another step\n");
7      }
8      fprintf(stderr, "Hare done running!\n");
9  }
```

Unlike `pthread_join()` which allows the target thread to pass back a value to the calling thread using a second parameter, `uv_thread_join()` does not. To send values use *Inter-thread communication*.

5.2 Synchronization Primitives

This section is purposely spartan. This book is not about threads, so I only catalogue any surprises in the libuv APIs here. For the rest you can look at the pthreads man pages.

5.2.1 Mutexes

The mutex functions are a **direct** map to the pthread equivalents.

libuv mutex functions

```
UV_EXTERN int uv_mutex_init(uv_mutex_t* handle);
UV_EXTERN void uv_mutex_destroy(uv_mutex_t* handle);
UV_EXTERN void uv_mutex_lock(uv_mutex_t* handle);
UV_EXTERN int uv_mutex_trylock(uv_mutex_t* handle);
UV_EXTERN void uv_mutex_unlock(uv_mutex_t* handle);
```

The `uv_mutex_init()` and `uv_mutex_trylock()` functions will return 0 on success, -1 on error instead of error codes.

If *libuv* has been compiled with debugging enabled, `uv_mutex_destroy()`, `uv_mutex_lock()` and `uv_mutex_unlock()` will `abort()` on error. Similarly `uv_mutex_trylock()` will abort if the error is anything *other than* `EAGAIN`.

Recursive mutexes are supported by some platforms, but you should not rely on them. The BSD mutex implementation will raise an error if a thread which has locked a mutex attempts to lock it again. For example, a construct like:

```
uv_mutex_lock(a_mutex);
uv_thread_create(thread_id, entry, (void *)a_mutex);
uv_mutex_lock(a_mutex);
// more things here
```

can be used to wait until another thread initializes some stuff and then unlocks a_mutex but will lead to your program crashing if in debug mode, or return an error in the second call to uv_mutex_lock().

Note: Mutexes on linux support attributes for a recursive mutex, but the API is not exposed via libuv.

5.2.2 Locks

Read-write locks are the other synchronization primitive supported. TODO some DB read/write example

5.2.3 Others

Semaphores and condition variables are not implemented yet. Their are a couple of patches for condition variable support in libuv [1] [2], but since the Windows condition variable system is available only from Windows Vista and Windows Server 2008 onwards [3], the patches haven't been merged.

5.3 libuv work queue

uv_queue_work() is a convenience function that allows an application to run a task in a separate thread, and have a callback that is triggered when the task is done. A seemingly simple function, what makes uv_queue_work() tempting is that it allows potentially any third-party libraries to be used with the event-loop paradigm. When you use event loops, it is *imperative to make sure that no function which runs periodically in the loop thread blocks when performing I/O or is a serious CPU hog*, because this means the loop slows down and events are not being dealt with at full capacity.

But a lot of existing code out there features blocking functions (for example a routine which performs I/O under the hood) to be used with threads if you want responsiveness (the classic 'one thread per client' server model), and getting them to play with an event loop library generally involves rolling your own system of running the task in a separate thread. libuv just provides a convenient abstraction for this.

Here is a simple example inspired by node.js is cancer. We are going to calculate fibonacci numbers, sleeping a bit along the way, but run it in a separate thread so that the blocking and CPU bound task does not prevent the event loop from performing other activities.

queue-work/main.c - lazy fibonacci

```
1  void fib(uv_work_t *req) {
2      int n = *(int *) req->data;
3      if (random() % 2)
4          sleep(1);
5      else
6          sleep(3);
7      long fib = fib_(n);
```

---

[1] https://github.com/nikhilm/libuv/compare/condvar
[2] https://github.com/bnoordhuis/libuv/compare/uv_cond
[3] http://msdn.microsoft.com/en-us/library/windows/desktop/ms683469(v=vs.85).aspx

---

```
8          fprintf(stderr, "%dth fibonacci is %lu\n", n, fib);
9    }
10
11   void after_fib(uv_work_t *req) {
12          fprintf(stderr, "Done calculating %dth fibonacci\n", *(int *) req->data);
13   }
```

The actual task function is simple, nothing to show that it is going to be run in a separate thread. The `uv_work_t` structure is the clue. You can pass arbitrary data through it using the `void* data` field and use it to communicate to and from the thread. But be sure you are using proper locks if you are changing things while both threads may be running.

The trigger is `uv_queue_work`:

queue-work/main.c

```
1    int main() {
2          loop = uv_default_loop();
3
4          int data[FIB_UNTIL];
5          uv_work_t req[FIB_UNTIL];
6          int i;
7          for (i = 0; i < FIB_UNTIL; i++) {
8                data[i] = i;
9                req[i].data = (void *) &data[i];
10               uv_queue_work(loop, &req[i], fib, after_fib);
11         }
12
13         return uv_run(loop);
14   }
```

The thread function will be launched in a separate thread, passed the `uv_work_t` structure and once the function returns, the *after* function will be called, again with the same structure.

For writing wrappers to blocking libraries, a common *pattern* is to use a baton to exchange data.

5.4 Inter-thread communication

Sometimes you want various threads to actually send each other messages *while* they are running. For example you might be running some long duration task in a separate thread (perhaps using `uv_queue_work`) but want to notify progress to the main thread. This is a simple example of having a download manager informing the user of the status of running downloads.

progress/main.c

```
1    uv_loop_t *loop;
2    uv_async_t async;
3
4    int main() {
5          loop = uv_default_loop();
6
7          uv_work_t req;
8          int size = 10240;
```

```
9       req.data = (void*) &size;
10
11      uv_async_init(loop, &async, print_progress);
12      uv_queue_work(loop, &req, fake_download, after);
13
14      return uv_run(loop);
15   }
```

The async thread communication works *on loops* so although any thread can be the message sender, only threads with libuv loops can be receivers (or rather the loop is the receiver). libuv will invoke the callback (`print_progress`) with the async watcher whenever it receives a message.

> **Warning:** It is important to realize that the message send is *async*, the callback may be invoked immediately after `uv_async_send` is called in another thread, or it may be invoked after some time. libuv may also combine multiple calls to `uv_async_send` and invoke your callback only once. The only guarantee that libuv makes is – The callback function is called *at least once* after the call to `uv_async_send`. If you have no pending calls to `uv_async_send`, the callback won't be called. If you make two or more calls, and libuv hasn't had a chance to run the callback yet, it *may* invoke your callback *only once* for the multiple invocations of `uv_async_send`. Your callback will never be called twice for just one event.

progress/main.c

```
1   void fake_download(uv_work_t *req) {
2       int size = *((int*) req->data);
3       int downloaded = 0;
4       double percentage;
5       while (downloaded < size) {
6           percentage = downloaded*100.0/size;
7           async.data = (void*) &percentage;
8           uv_async_send(&async);
9
10          sleep(1);
11          downloaded += (200+random())%1000; // can only download max 1000bytes/sec,
12                                              // but at least a 200;
13      }
14   }
```

In the download function we modify the progress indicator and queue the message for delivery with `uv_async_send`. Remember: `uv_async_send` is also non-blocking and will return immediately.

progress/main.c

```
1   void print_progress(uv_async_t *handle, int status /*UNUSED*/) {
2       double percentage = *((double*) handle->data);
3       fprintf(stderr, "Downloaded %.2f%%\n", percentage);
4   }
```

The callback is a standard libuv pattern, extracting the data from the watcher.

Finally it is important to remember to clean up the watcher.

progress/main.c

```
1  void after(uv_work_t *req) {
2      fprintf(stderr, "Download complete\n");
3      uv_close((uv_handle_t*) &async, NULL);
4  }
```

After this example, which showed the abuse of the `data` field, bnoordhuis pointed out that using the `data` field is not thread safe, and `uv_async_send()` is actually only meant to wake up the event loop. Use a mutex or rwlock to ensure accesses are performed in the right order.

> **Warning:** mutexes and rwlocks **DO NOT** work inside a signal handler, whereas `uv_async_send` does.

One use case where `uv_async_send` is required is when interoperating with libraries that require thread affinity for their functionality. For example in node.js, a v8 engine instance, contexts and its objects are bound to the thread that the v8 instance was started in. Interacting with v8 data structures from another thread can lead to undefined results. Now consider some node.js module which binds a third party library. It may go something like this:

1. In node, the third party library is set up with a JavaScript callback to be invoked for more information:

   ```
   var lib = require('lib');
   lib.on_progress(function() {
       console.log("Progress");
   });

   lib.do();

   // do other stuff
   ```

2. `lib.do` is supposed to be non-blocking but the third party lib is blocking, so the binding uses `uv_queue_work`.

3. The actual work being done in a separate thread wants to invoke the progress callback, but cannot directly call into v8 to interact with JavaScript. So it uses `uv_async_send`.

4. The async callback, invoked in the main loop thread, which is the v8 thread, then interacts with v8 to invoke the JavaScript callback.

PROCESSES

libuv offers considerable child process management, abstracting the platform differences and allowing communication with the child process using streams or named pipes.

A common idiom in Unix is for every process to do one thing and do it well. In such a case, a process often uses multiple child processes to achieve tasks (similar to using pipes in shells). A multi-process model with messages may also be easier to reason about compared to one with threads and shared memory.

A common refrain against event-based programs is that they cannot take advantage of multiple cores in modern computers. In a multi-threaded program the kernel can perform scheduling and assign different threads to different cores, improving performance. But an event loop has only one thread. The workaround can be to launch multiple processes instead, with each process running an event loop, and each process getting assigned to a separate CPU core.

6.1 Spawning child processes

The simplest case is when you simply want to launch a process and know when it exits. This is achieved using `uv_spawn`.

spawn/main.c

```
1  uv_loop_t *loop;
2  uv_process_t child_req;
3  uv_process_options_t options;
4
5  int main() {
6      loop = uv_default_loop();
7
8      char* args[3];
9      args[0] = "mkdir";
10     args[1] = "test-dir";
11     args[2] = NULL;
12
13     options.exit_cb = on_exit;
14     options.file = "mkdir";
15     options.args = args;
16
17     if (uv_spawn(loop, &child_req, options)) {
18         fprintf(stderr, "%s\n", uv_strerror(uv_last_error(loop)));
19         return 1;
20     }
21
```

```
22        return uv_run(loop);
23    }
```

The `uv_process_t` struct only acts as the watcher, all options are set via `uv_process_options_t`. To simply launch a process, you need to set only the `file` and `args` fields. `file` is the program to execute. Since `uv_spawn` uses execvp internally, there is no need to supply the full path. Finally as per underlying conventions, **the arguments array has to be one larger than the number of arguments, with the last element being NULL**.

After the call to `uv_spawn`, `uv_process_t.pid` will contain the process ID of the child process.

The exit callback will be invoked with the *exit status* and the type of *signal* which caused the exit.

spawn/main.c

```
1  void on_exit(uv_process_t *req, int exit_status, int term_signal) {
2      fprintf(stderr, "Process exited with status %d, signal %d\n", exit_status, term_signal);
3      uv_close((uv_handle_t*) req, NULL);
4  }
```

It is **required** to close the process watcher after the process exits.

6.2 Changing process parameters

Before the child process is launched you can control the execution environment using fields in `uv_process_options_t`.

6.2.1 Change execution directory

Set `uv_process_options_t.cwd` to the corresponding directory.

6.2.2 Set environment variables

`uv_process_options_t.env` is an array of strings, each of the form `VAR=VALUE` used to set up the environment variables for the process. Set this to `NULL` to inherit the environment from the parent (this) process.

6.2.3 Option flags

Setting `uv_process_options_t.flags` to a bitwise OR of the following flags, modifies the child process behaviour:

- `UV_PROCESS_SETUID` - sets the child's execution user ID to `uv_process_options_t.uid`.
- `UV_PROCESS_SETGID` - sets the child's execution group ID to `uv_process_options_t.gid`.

Changing the UID/GID is only supported on Unix, `uv_spawn` will fail on Windows with `UV_ENOTSUP`.

- `UV_PROCESS_WINDOWS_VERBATIM_ARGUMENTS` - No quoting or escaping of `uv_process_options_t.args` is done on Windows. Ignored on Unix.
- `UV_PROCESS_DETACHED` - Starts the child process in a new session, which will keep running after the parent process exits. See example below.

6.3 Detaching processes

Passing the flag `UV_PROCESS_DETACHED` can be used to launch daemons, or child processes which are independent of the parent so that the parent exiting does not affect it.

detach/main.c

```c
1  int main() {
2      loop = uv_default_loop();
3
4      char* args[3];
5      args[0] = "sleep";
6      args[1] = "100";
7      args[2] = NULL;
8
9      options.exit_cb = NULL;
10     options.file = "sleep";
11     options.args = args;
12     options.flags = UV_PROCESS_DETACHED;
13
14     if (uv_spawn(loop, &child_req, options)) {
15         fprintf(stderr, "%s\n", uv_strerror(uv_last_error(loop)));
16         return 1;
17     }
18     fprintf(stderr, "Launched sleep with PID %d\n", child_req.pid);
19     uv_unref((uv_handle_t*) &child_req);
20
21     return uv_run(loop);
22 }
```

Just remember that the watcher is still monitoring the child, so your program won't exit. Use `uv_unref()` if you want to be more *fire-and-forget*.

6.4 Signals and termination

libuv wraps the standard `kill(2)` system call on Unix and implements one with similar semantics on Windows, with *one caveat*: `uv_kill` on Windows only supports `SIGTERM`, `SIGINT` and `SIGKILL`, all of which lead to termination of the process. The signature of `uv_kill` is:

```c
uv_err_t uv_kill(int pid, int signum);
```

For processes started using libuv, you may use `uv_process_kill` instead, which accepts the `uv_process_t` watcher as the first argument, rather than the pid. In this case, **remember to call** `uv_close` on the watcher.

6.5 Child Process I/O

A normal, newly spawned process has its own set of file descriptors, with 0, 1 and 2 being `stdin`, `stdout` and `stderr` respectively. Sometimes you may want to share file descriptors with the child. For example, perhaps your applications launches a sub-command and you want any errors to go in the log file, but ignore `stdout`. For this you'd like to have `stderr` of the child to be displayed. In this case, libuv supports *inheriting* file descriptors. In this sample, we invoke the test program, which is:

proc-streams/test.c

```c
#include <stdio.h>

int main()
{
    fprintf(stderr, "This is stderr\n");
    printf("This is stdout\n");
    return 0;
}
```

The actual program `proc-streams` runs this while inheriting only `stderr`. The file descriptors of the child process are set using the `stdio` field in `uv_process_options_t`. First set the `stdio_count` field to the number of file descriptors being set. `uv_process_options_t.stdio` is an array of `uv_stdio_container_t`, which is:

```c
typedef struct uv_stdio_container_s {
  uv_stdio_flags flags;

  union {
    uv_stream_t* stream;
    int fd;
  } data;
} uv_stdio_container_t;
```

where flags can have several values. Use `UV_IGNORE` if it isn't going to be used. If the first three `stdio` fields are marked as `UV_IGNORE` they'll redirect to `/dev/null`.

Since we want to pass on an existing descriptor, we'll use `UV_INHERIT_FD`. Then we set the `fd` to `stderr`.

proc-streams/main.c

```c
1   int main() {
2       loop = uv_default_loop();
3
4       /* ... */
5
6       options.stdio_count = 3;
7       uv_stdio_container_t child_stdio[3];
8       child_stdio[0].flags = UV_IGNORE;
9       child_stdio[1].flags = UV_IGNORE;
10      child_stdio[2].flags = UV_INHERIT_FD;
11      child_stdio[2].data.fd = 2;
12      options.stdio = child_stdio;
13
14      options.exit_cb = on_exit;
15      options.file = args[0];
16      options.args = args;
17
18
19      if (uv_spawn(loop, &child_req, options)) {
20          fprintf(stderr, "%s\n", uv_strerror(uv_last_error(loop)));
21          return 1;
22      }
23
24      return uv_run(loop);
25  }
```

If you run `proc-stream` you'll see that only the line "This is stderr" will be displayed. Try marking `stdout` as being inherited and see the output.

It is dead simple to apply this redirection to streams. By setting `flags` to `UV_INHERIT_STREAM` and setting `data.stream` to the stream in the parent process, the child process can treat that stream as standard I/O. This can be used to implement something like CGI.

A sample CGI script/executable is:

cgi/tick.c

```c
#include <stdio.h>
#include <unistd.h>

int main() {
    int i;
    for (i = 0; i < 10; i++) {
        printf("tick\n");
        fflush(stdout);
        sleep(1);
    }
    printf("BOOM!\n");
    return 0;
}
```

The CGI server combines the concepts from this chapter and *Networking* so that every client is sent ten ticks after which that connection is closed.

cgi/main.c

```c
void on_new_connection(uv_stream_t *server, int status) {
    uv_tcp_t *client = (uv_tcp_t*) malloc(sizeof(uv_tcp_t));
    uv_tcp_init(loop, client);
    if (uv_accept(server, (uv_stream_t*) client) == 0) {
        invoke_cgi_script(client);
    }
    else {
        uv_close((uv_handle_t*) client, NULL);
    }
}
```

Here we simply accept the TCP connection and pass on the socket (*stream*) to `invoke_cgi_script`.

cgi/main.c

```c
void invoke_cgi_script(uv_tcp_t *client) {

    /* ... finding the executable path and setting up arguments ... */

    options.stdio_count = 3;
    uv_stdio_container_t child_stdio[3];
    child_stdio[0].flags = UV_IGNORE;
    child_stdio[1].flags = UV_INHERIT_STREAM;
    child_stdio[1].data.stream = (uv_stream_t*) client;
    child_stdio[2].flags = UV_IGNORE;
```

```
11      options.stdio = child_stdio;
12
13      options.exit_cb = on_exit;
14      options.file = args[0];
15      options.args = args;
16
17      child_req.data = (void*) client;
18      if (uv_spawn(loop, &child_req, options)) {
19          fprintf(stderr, "%s\n", uv_strerror(uv_last_error(loop)));
20          return;
21      }
22  }
```

The `stdout` of the CGI script is set to the socket so that whatever our tick script prints, gets sent to the client. By using processes, we can offload the read/write buffering to the operating system, so in terms of convenience this is great. Just be warned that creating processes is a costly task.

6.6 Pipes

libuv's `uv_pipe_t` structure is slightly confusing to Unix programmers, because it immediately conjures up | and pipe(7). But `uv_pipe_t` is not related to anonymous pipes, rather it has two uses:

1. Stream API - It acts as the concrete implementation of the `uv_stream_t` API for providing a FIFO, streaming interface to local file I/O. This is performed using `uv_pipe_open` as covered in *Buffers and Streams*. You could also use it for TCP/UDP, but there are already convenience functions and structures for them.

2. IPC mechanism - `uv_pipe_t` can be backed by a Unix Domain Socket or Windows Named Pipe to allow multiple processes to communicate. This is discussed below.

6.6.1 Parent-child IPC

A parent and child can have one or two way communication over a pipe created by settings `uv_stdio_container_t.flags` to a bit-wise combination of `UV_CREATE_PIPE` and `UV_READABLE_PIPE` or `UV_WRITABLE_PIPE`. The read/write flag is from the perspective of the child process.

6.6.2 Arbitrary process IPC

Since domain sockets [1] can have a well known name and a location in the file-system they can be used for IPC between unrelated processes. The D-BUS system used by open source desktop environments uses domain sockets for event notification. Various applications can then react when a contact comes online or new hardware is detected. The MySQL server also runs a domain socket on which clients can interact with it.

When using domain sockets, a client-server pattern is usually followed with the creator/owner of the socket acting as the server. After the initial setup, messaging is no different from TCP, so we'll re-use the echo server example.

pipe-echo-server/main.c

```
1  int main() {
2      loop = uv_default_loop();
3
```

---

[1] In this section domain sockets stands in for named pipes on Windows as well.

```
4        uv_pipe_t server;
5        uv_pipe_init(loop, &server, 0);
6
7        signal(SIGINT, remove_sock);
8
9        if (uv_pipe_bind(&server, "echo.sock")) {
10           fprintf(stderr, "Bind error %s\n", uv_err_name(uv_last_error(loop)));
11           return 1;
12       }
13       if (uv_listen((uv_stream_t*) &server, 128, on_new_connection)) {
14           fprintf(stderr, "Listen error %s\n", uv_err_name(uv_last_error(loop)));
15           return 2;
16       }
17       return uv_run(loop);
18   }
```

We name the socket `echo.sock` which means it will be created in the local directory. This socket now behaves no different from TCP sockets as far as the stream API is concerned. You can test this server using netcat:

```
$ nc -U /path/to/echo.sock
```

A client which wants to connect to a domain socket will use:

```
void uv_pipe_connect(uv_connect_t *req, uv_pipe_t *handle, const char *name, uv_connect_cb cb);
```

where `name` will be `echo.sock` or similar.

6.6.3 Sending file descriptors over pipes

The cool thing about domain sockets is that file descriptors can be exchanged between processes by sending them over a domain socket. This allows processes to hand off their I/O to other processes. Applications include load-balancing servers, worker processes and other ways to make optimum use of CPU.

> **Warning:** On Windows, only file descriptors representing TCP sockets can be passed around.

To demonstrate, we will look at a echo server implementation that hands of clients to worker processes in a round-robin fashion. This program is a bit involved, and while only snippets are included in the book, it is recommended to read the full code to really understand it.

The worker process is quite simple, since the file-descriptor is handed over to it by the master.

multi-echo-server/worker.c

```
1  uv_loop_t *loop;
2  uv_pipe_t queue;
3
4  int main() {
5      loop = uv_default_loop();
6
7      uv_pipe_init(loop, &queue, 1);
8      uv_pipe_open(&queue, 0);
9      uv_read2_start(&queue, alloc_buffer, on_new_connection);
10     return uv_run(loop);
11 }
```

queue is the pipe connected to the master process on the other end, along which new file descriptors get sent. We use the `read2` function to express interest in file descriptors. It is important to set the `ipc` argument of `uv_pipe_init` to 1 to indicate this pipe will be used for inter-process communication! Since the master will write the file handle to the standard input of the worker, we connect the pipe to `stdin` using `uv_pipe_open`.

multi-echo-server/worker.c

```c
void on_new_connection(uv_pipe_t *q, ssize_t nread, uv_buf_t buf, uv_handle_type pending) {
    if (pending == UV_UNKNOWN_HANDLE) {
        // error!
        return;
    }

    uv_pipe_t *client = (uv_pipe_t*) malloc(sizeof(uv_pipe_t));
    uv_pipe_init(loop, client, 0);
    if (uv_accept((uv_stream_t*) q, (uv_stream_t*) client) == 0) {
        fprintf(stderr, "Worker %d: Accepted fd %d\n", getpid(), client->fd);
        uv_read_start((uv_stream_t*) client, alloc_buffer, echo_read);
    }
    else {
        uv_close((uv_handle_t*) client, NULL);
    }
}
```

Although `accept` seems odd in this code, it actually makes sense. What `accept` traditionally does is get a file descriptor (the client) from another file descriptor (The listening socket). Which is exactly what we do here. Fetch the file descriptor (`client`) from queue. From this point the worker does standard echo server stuff.

Turning now to the master, let's take a look at how the workers are launched to allow load balancing.

multi-echo-server/main.c

```c
uv_loop_t *loop;

struct child_worker {
    uv_process_t req;
    uv_process_options_t options;
    uv_pipe_t pipe;
} *workers;
```

The `child_worker` structure wraps the process, and the pipe between the master and the individual process.

multi-echo-server/main.c

```c
void setup_workers() {
    // ...

    // launch same number of workers as number of CPUs
    uv_cpu_info_t *info;
    int cpu_count;
    uv_cpu_info(&info, &cpu_count);
    uv_free_cpu_info(info, cpu_count);

    child_worker_count = cpu_count;
```

```
11
12        workers = calloc(sizeof(struct child_worker), cpu_count);
13        while (cpu_count--) {
14            struct child_worker *worker = &workers[cpu_count];
15            uv_pipe_init(loop, &worker->pipe, 1);
16
17            uv_stdio_container_t child_stdio[3];
18            child_stdio[0].flags = UV_CREATE_PIPE | UV_READABLE_PIPE;
19            child_stdio[0].data.stream = (uv_stream_t*) &worker->pipe;
20            child_stdio[1].flags = UV_IGNORE;
21            child_stdio[2].flags = UV_INHERIT_FD;
22            child_stdio[2].data.fd = 2;
23
24            worker->options.stdio = child_stdio;
25            worker->options.stdio_count = 3;
26
27            worker->options.exit_cb = on_exit;
28            worker->options.file = args[0];
29            worker->options.args = args;
30
31            uv_spawn(loop, &worker->req, worker->options);
32            fprintf(stderr, "Started worker %d\n", worker->req.pid);
33        }
34    }
```

In setting up the workers, we use the nifty libuv function `uv_cpu_info` to get the number of CPUs so we can launch an equal number of workers. Again it is important to initialize the pipe acting as the IPC channel with the third argument as 1. We then indicate that the child process' `stdin` is to be a readable pipe (from the point of view of the child). Everything is straightforward till here. The workers are launched and waiting for file descriptors to be written to their pipes.

It is in `on_new_connection` (the TCP infrastructure is initialized in `main()`), that we accept the client socket and pass it along to the next worker in the round-robin.

multi-echo-server/main.c

```
1   void on_new_connection(uv_stream_t *server, int status) {
2       if (status == -1) {
3           // error!
4           return;
5       }
6
7       uv_pipe_t *client = (uv_pipe_t*) malloc(sizeof(uv_pipe_t));
8       uv_pipe_init(loop, client, 0);
9       if (uv_accept(server, (uv_stream_t*) client) == 0) {
10          uv_write_t *write_req = (uv_write_t*) malloc(sizeof(uv_write_t));
11          dummy_buf = uv_buf_init(".", 1);
12          struct child_worker *worker = &workers[round_robin_counter];
13          uv_write2(write_req, (uv_stream_t*) &worker->pipe, &dummy_buf, 1, (uv_stream_t*) client, NULL
14          round_robin_counter = (round_robin_counter + 1) % child_worker_count;
15      }
16      else {
17          uv_close((uv_handle_t*) client, NULL);
18      }
19  }
```

Again, the `uv_write2` call handles all the abstraction and it is simply a matter of passing in the file descriptor as the

right argument. With this our multi-process echo server is operational.

TODO what do the write2/read2 functions do with the buffers?

MULTIPLE EVENT LOOPS

UTILITIES

This chapter catalogues tools and techniques which are useful for common tasks. The libev man page already covers some patterns which can be adopted to libuv through simple API changes. It also covers parts of the libuv API that don't require entire chapters dedicated to them.

8.1 Timers

Timers invoke the callback after a certain time has elapsed since the timer was started. libuv timers can also be set to invoke at regular intervals instead of just once.

Simple use is to init a watcher and start it with a `timeout`, and optional `repeat`. Timers can be stopped at any time.

```
uv_timer_t timer_req;

uv_timer_init(loop, &timer_req);
uv_timer_start(&timer_req, callback, 5000, 2000);
```

will start a repeating timer, which first starts 5 seconds (the `timeout`) after the execution of `uv_timer_start`, then repeats every 2 seconds (the `repeat`). Use:

```
uv_timer_stop(&timer_req);
```

to stop the timer. This can be used safely from within the callback as well.

The repeat interval can be modified at any time with:

```
uv_timer_set_repeat(uv_timer_t *timer, int64_t repeat);
```

which will take effect **when possible**. If this function is called from a timer callback, it means:

- If the timer was non-repeating, the timer has already been stopped. Use `uv_timer_start` again.

- If the timer is repeating, the next timeout has already been scheduled, so the old repeat interval will be used once more before the timer switches to the new interval.

The utility function:

```
int uv_timer_again(uv_timer_t *)
```

applies **only to repeating timers** and is equivalent to stopping the timer and then starting it with both initial `timeout` and `repeat` set to the old `repeat` value. If the timer hasn't been started it fails (error code `UV_EINVAL`) and returns -1.

An actual timer example is in the *reference count section*.

8.2 Event loop reference count

The event loop only runs as long as there are active watchers. This system works by having every watcher increase the reference count of the event loop when it is started and decreasing the reference count when stopped. It is also possible to manually change the reference count of handles using:

```
void uv_ref(uv_handle_t*);
void uv_unref(uv_handle_t*);
```

These functions can be used to allow a loop to exit even when a watcher is active or to use custom objects to keep the loop alive.

The former can be used with interval timers. You might have a garbage collector which runs every X seconds, or your network service might send a heartbeat to others periodically, but you don't want to have to stop them along all clean exit paths or error scenarios. Or you want the program to exit when all your other watchers are done. In that case just unref the timer immediately after creation so that if it is the only watcher running then `uv_run` will still exit.

The later is used in node.js where some libuv methods are being bubbled up to the JS API. A `uv_handle_t` (the superclass of all watchers) is created per JS object and can be ref/unrefed.

ref-timer/main.c

```
1   uv_loop_t *loop;
2   uv_timer_t gc_req;
3   uv_timer_t fake_job_req;
4
5   int main() {
6       loop = uv_default_loop();
7
8       uv_timer_init(loop, &gc_req);
9       uv_unref((uv_handle_t*) &gc_req);
10
11      uv_timer_start(&gc_req, gc, 0, 2000);
12
13      // could actually be a TCP download or something
14      uv_timer_init(loop, &fake_job_req);
15      uv_timer_start(&fake_job_req, fake_job, 9000, 0);
16      return uv_run(loop);
17  }
```

We initialize the garbage collector timer, then immediately `unref` it. Observe how after 9 seconds, when the fake job is done, the program automatically exits, even though the garbage collector is still running.

8.3 Idle watcher pattern

The callbacks of idle watchers are only invoked when the event loop has no other pending events. In such a situation they are invoked once every iteration of the loop. The idle callback can be used to perform some very low priority activity. For example, you could dispatch a summary of the daily application performance to the developers for analysis during periods of idleness, or use the application's CPU time to perform SETI calculations :) An idle watcher is also useful in a GUI application. Say you are using an event loop for a file download. If the TCP socket is still being established and no other events are present your event loop will pause (**block**), which means your progress bar will freeze and the user will think the application crashed. In such a case queue up and idle watcher to keep the UI operational.

idle-compute/main.c

```
1  uv_loop_t *loop;
2  uv_fs_t stdin_watcher;
3  uv_idle_t idler;
4  char buffer[1024];
5
6  int main() {
7      loop = uv_default_loop();
8
9      uv_idle_init(loop, &idler);
10
11     uv_fs_read(loop, &stdin_watcher, 1, buffer, 1024, -1, on_type);
12     uv_idle_start(&idler, crunch_away);
13     return uv_run(loop);
14 }
```

Here we initialize the idle watcher and queue it up along with the actual events we are interested in. `crunch_away` will now be called repeatedly until the user types something and presses Return. Then it will be interrupted for a brief amount as the loop deals with the input data, after which it will keep calling the idle callback again.

idle-compute/main.c

```
1  void crunch_away(uv_idle_t* handle, int status) {
2      // Compute extra-terrestrial life
3      // fold proteins
4      // computer another digit of PI
5      // or similar
6      fprintf(stderr, "Computing PI...\n");
7      // just to avoid overwhelming your terminal emulator
8      uv_idle_stop(handle);
9  }
```

8.4 Passing data to worker thread

When using `uv_queue_work` you'll usually need to pass complex data through to the worker thread. The solution is to use a `struct` and set `uv_work_t.data` to point to it. A slight variation is to have the `uv_work_t` itself as the first member of this struct (called a baton [1]). This allows cleaning up the work request and all the data in one free call.

```
1  struct ftp_baton {
2      uv_work_t req;
3      char *host;
4      int port;
5      char *username;
6      char *password;
7  }
```

```
1  ftp_baton *baton = (ftp_baton*) malloc(sizeof(ftp_baton));
2  baton->req.data = (void*) baton;
3  baton->host = strdup("my.webhost.com");
4  baton->port = 21;
```

---

[1] mfp is My Fancy Plugin

```
5   // ...
6
7   uv_queue_work(loop, &baton->req, ftp_session, ftp_cleanup);
```

Here we create the baton and queue the task.

Now the task function can extract the data it needs:

```
1   void ftp_session(uv_work_t *req) {
2       ftp_baton *baton = (ftp_baton*) req->data;
3
4       fprintf(stderr, "Connecting to %s\n", baton->host);
5   }
6
7   void ftp_cleanup(uv_work_t *req) {
8       ftp_baton *baton = (ftp_baton*) req->data;
9
10      free(baton->host);
11      // ...
12      free(baton);
13  }
```

We then free the baton which also frees the watcher.

8.5 External I/O with polling

Usually third-party libraries will handle their own I/O, and keep track of their sockets and other files internally. In this case it isn't possible to use the standard stream I/O operations, but the library can still be integrated into the libuv event loop. All that is required is that the library allow you to access the underlying file descriptors and provide functions that process tasks in small increments as decided by your application. Some libraries though will not allow such access, providing only a standard blocking function which will perform the entire I/O transaction and only then return. It is unwise to use these in the event loop thread, use the *libuv work queue* instead. Of course this will also mean losing granular control on the library.

The `uv_poll` section of libuv simply watches file descriptors using the operating system notification mechanism. In some sense, all the I/O operations that libuv implements itself are also backed by `uv_poll` like code. Whenever the OS notices a change of state in file descriptors being polled, libuv will invoke the associated callback.

Here we will walk through a simple download manager that will use libcurl to download files. Rather than give all control to libcurl, we'll instead be using the libuv event loop, and use the non-blocking, async multi interface to progress with the download whenever libuv notifies of I/O readiness.

uvwget/main.c - The setup

```
1   #include <stdio.h>
2   #include <stdlib.h>
3   #include <sys/select.h>
4   #include <uv.h>
5   #include <curl/curl.h>
6
7   uv_loop_t *loop;
8   CURLM *curl_handle;
9   uv_timer_t timeout;
10
11  int main(int argc, char **argv) {
```

```
12      loop = uv_default_loop();
13
14      if (argc <= 1)
15          return 0;
16
17      if (curl_global_init(CURL_GLOBAL_ALL)) {
18          fprintf(stderr, "Could not init cURL\n");
19          return 1;
20      }
21
22      uv_timer_init(loop, &timeout);
23
24      curl_handle = curl_multi_init();
25      curl_multi_setopt(curl_handle, CURLMOPT_SOCKETFUNCTION, handle_socket);
26      curl_multi_setopt(curl_handle, CURLMOPT_TIMERFUNCTION, start_timeout);
27
28      while (argc-- > 1) {
29          add_download(argv[argc], argc);
30      }
31
32      uv_run(loop);
33      curl_multi_cleanup(curl_handle);
34      return 0;
35  }
```

The way each library is integrated with libuv will vary. In the case of libcurl, we can register two callbacks. The socket callback `handle_socket` is invoked whenever the state of a socket changes and we have to start polling it. `start_timeout` is called by libcurl to notify us of the next timeout interval, after which we should drive libcurl forward regardless of I/O status. This is so that libcurl can handle errors or do whatever else is required to get the download moving.

Our downloader is to be invoked as:

```
$ ./uvwget [url1] [url2] ...
```

So we add each argument as an URL

uvwget/main.c - Adding urls

```
1   void add_download(const char *url, int num) {
2       char filename[50];
3       sprintf(filename, "%d.download", num);
4       FILE *file;
5
6       file = fopen(filename, "w");
7       if (file == NULL) {
8           fprintf(stderr, "Error opening %s\n", filename);
9           return;
10      }
11
12      CURL *handle = curl_easy_init();
13      curl_easy_setopt(handle, CURLOPT_WRITEDATA, file);
14      curl_easy_setopt(handle, CURLOPT_URL, url);
15      curl_multi_add_handle(curl_handle, handle);
16      fprintf(stderr, "Added download %s -> %s\n", url, filename);
17  }
```

We let libcurl directly write the data to a file, but much more is possible if you so desire.

start_timeout will be called immediately the first time by libcurl, so things are set in motion. This simply starts a libuv timer which drives curl_multi_socket_action with CURL_SOCKET_TIMEOUT whenever it times out. curl_multi_socket_action is what drives libcurl, and what we call whenever sockets change state. But before we go into that, we need to poll on sockets whenever handle_socket is called.

uvwget/main.c - Setting up polling

```c
int handle_socket(CURL *easy, curl_socket_t s, int action, void *userp, void *socketp) {
    uv_poll_t *poll_fd;
    if (action == CURL_POLL_IN || action == CURL_POLL_OUT) {
        if (socketp) {
            poll_fd = (uv_poll_t*) socketp;
        }
        else {
            poll_fd = (uv_poll_t*) malloc(sizeof(uv_poll_t));
            uv_poll_init(loop, poll_fd, s);
        }
        curl_multi_assign(curl_handle, s, (void *) poll_fd);
    }

    switch (action) {
        case CURL_POLL_IN:
            uv_poll_start(poll_fd, UV_READABLE, curl_perform);
            break;
        case CURL_POLL_OUT:
            uv_poll_start(poll_fd, UV_WRITABLE, curl_perform);
            break;
        case CURL_POLL_REMOVE:
            if (socketp) {
                uv_poll_stop((uv_poll_t*) socketp);
                uv_close((uv_handle_t*) socketp, (uv_close_cb) free);
                curl_multi_assign(curl_handle, s, NULL);
            }
            break;
        default:
            abort();
    }

    return 0;
}
```

We are interested in the socket fd s, and the action. For every socket we create a uv_poll_t handle if it doesn't exist, and associate it with the socket using curl_multi_assign. This way socketp points to it whenever the callback is invoked.

In the case that the download is done or fails, libcurl requests removal of the poll. So we stop and free the poll handle.

Depending on what events libcurl wishes to watch for, we start polling with UV_READABLE or UV_WRITABLE. Now libuv will invoke the poll callback whenever the socket is ready for reading or writing. Calling uv_poll_start multiple times on the same handle is acceptable, it will just update the events mask with the new value. curl_perform is the crux of this program.

uvwget/main.c - Setting up polling

```c
void curl_perform(uv_poll_t *req, int status, int events) {
    uv_timer_stop(&timeout);
    int running_handles;
    int flags = 0;
    if (events & UV_READABLE) flags |= CURL_CSELECT_IN;
    if (events & UV_WRITABLE) flags |= CURL_CSELECT_OUT;

    curl_multi_socket_action(curl_handle, req->fd, flags, &running_handles);

    char *done_url;

    CURLMsg *message;
    int pending;
    while ((message = curl_multi_info_read(curl_handle, &pending))) {
        switch (message->msg) {
        case CURLMSG_DONE:
            curl_easy_getinfo(message->easy_handle, CURLINFO_EFFECTIVE_URL, &done_url);
            printf("%s DONE\n", done_url);

            curl_multi_remove_handle(curl_handle, message->easy_handle);
            curl_easy_cleanup(message->easy_handle);

            break;
        default:
            fprintf(stderr, "CURLMSG default\n");
            abort();
        }
    }
}
```

The first thing we do is to stop the timer, since there has been some progress in the interval. Then depending on what event triggered the callback, we inform libcurl of the same. Then we call `curl_multi_socket_action` with the socket that progressed and the flags informing about what events happened. At this point libcurl does all of its internal tasks in small increments, and will attempt to return as fast as possible, which is exactly what an evented program wants in its main thread. libcurl keeps queueing messages into its own queue about transfer progress. In our case we are only interested in transfers that are completed. So we extract these messages, and clean up handles whose transfers are done.

8.6 Check & Prepare watchers

TODO

8.7 Loading libraries

libuv provides a cross platform API to dynamically load shared libraries. This can be used to implement your own plugin/extension/module system and is used by node.js to implement `require()` support for bindings. The usage is quite simple as long as your library exports the right symbols. Be careful with sanity and security checks when loading third party code, otherwise your program will behave unpredicatably. This example implements a very simple plugin system which does nothing except print the name of the plugin.

Let us first look at the interface provided to plugin authors.

plugin/plugin.h

```
1  #ifndef UVBOOK_PLUGIN_SYSTEM
2  #define UVBOOK_PLUGIN_SYSTEM
3
4  void mfp_register(const char *name);
5
6  #endif
```

plugin/plugin.c

```
1  #include <stdio.h>
2
3  void mfp_register(const char *name) {
4      fprintf(stderr, "Registered plugin \"%s\"\n", name);
5  }
```

You can similarly add more functions that plugin authors can use to do useful things in your application [2]. A sample plugin using this API is:

plugin/hello.c

```
1  #include "plugin.h"
2
3  void initialize() {
4      mfp_register("Hello World!");
5  }
```

Our interface defines that all plugins should have an `initialize` function which will be called by the application. This plugin is compiled as a shared library and can be loaded by running our application:

```
$ ./plugin libhello.dylib
Loading libhello.dylib
Registered plugin "Hello World!"
```

This is done by using `uv_dlopen` to first load the shared library `libhello.dylib`. Then we get access to the `initialize` function using `uv_dlsym` and invoke it.

plugin/main.c

```
1  #include "plugin.h"
2
3  typedef void (*init_plugin_function)();
4
5  int main(int argc, char **argv) {
6      if (argc == 1) {
7          fprintf(stderr, "Usage: %s [plugin1] [plugin2] ...\n", argv[0]);
8          return 0;
9      }
10
11     uv_lib_t *lib = (uv_lib_t*) malloc(sizeof(uv_lib_t));
```

---

[2] I was first introduced to the term baton in this context, in Konstantin Käfer's excellent slides on writing node.js bindings – http://kkaefer.github.com/node-cpp-modules/#baton

```
12      while (--argc) {
13          fprintf(stderr, "Loading %s\n", argv[argc]);
14          if (uv_dlopen(argv[argc], lib)) {
15              fprintf(stderr, "Error: %s\n", uv_dlerror(lib));
16              continue;
17          }
18
19          init_plugin_function init_plugin;
20          if (uv_dlsym(lib, "initialize", (void **) &init_plugin)) {
21              fprintf(stderr, "dlsym error: %s\n", uv_dlerror(lib));
22              continue;
23          }
24
25          init_plugin();
26      }
27
28      return 0;
29  }
```

uv_dlopen expects a path to the shared library and sets the opaque uv_lib_t pointer. It returns 0 on success, -1 on error. Use uv_dlerror to get the error message.

uv_dlsym stores a pointer to the symbol in the second argument in the third argument. init_plugin_function is a function pointer to the sort of function we are looking for in the application's plugins.

8.8 TTY

Text terminals have supported basic formatting for a long time, with a pretty standardised command set. This formatting is often used by programs to improve the readability of terminal output. For example grep --colour. libuv provides the uv_tty_t abstraction (a stream) and related functions to implement the ANSI escape codes across all platforms. By this I mean that libuv converts ANSI codes to the Windows equivalent, and provides functions to get terminal information.

The first thing to do is to initialize a uv_tty_t with the file descriptor it reads/writes from. This is achieved with:

int uv_tty_init(uv_loop_t*, uv_tty_t*, uv_file fd, **int** readable)

If readable is false, uv_write calls to this stream will be **blocking**.

It is then best to use uv_tty_set_mode to set the mode to *normal* (0) which enables most TTY formatting, flow-control and other settings. *raw* mode (1) is also supported.

Remember to call uv_tty_reset_mode when your program exits to restore the state of the terminal. Just good manners. Another set of good manners is to be aware of redirection. If the user redirects the output of your command to a file, control sequences should not be written as they impede readability and grep. To check if the file descriptor is indeed a TTY, call uv_guess_handle with the file descriptor and compare the return value with UV_TTY.

Here is a simple example which prints white text on a red background:

tty/main.c

```
1  #include <stdio.h>
2  #include <string.h>
3  #include <unistd.h>
4  #include <uv.h>
5
```

```
6    uv_loop_t *loop;
7    uv_tty_t tty;
8    int main() {
9        loop = uv_default_loop();
10
11       uv_tty_init(loop, &tty, 1, 0);
12       uv_tty_set_mode(&tty, 0);
13
14       if (uv_guess_handle(1) == UV_TTY) {
15           uv_write_t req;
16           uv_buf_t buf;
17           buf.base = "\033[41;37m";
18           buf.len = strlen(buf.base);
19           uv_write(&req, (uv_stream_t*) &tty, &buf, 1, NULL);
20       }
21
22       uv_write_t req;
23       uv_buf_t buf;
24       buf.base = "Hello TTY\n";
25       buf.len = strlen(buf.base);
26       uv_write(&req, (uv_stream_t*) &tty, &buf, 1, NULL);
27       uv_tty_reset_mode();
28       return uv_run(loop);
29   }
```

The final TTY helper is `uv_tty_get_winsize()` which is used to get the width and height of the terminal and returns 0 on success. Here is a small program which does some animation using the function and character position escape codes.

tty-gravity/main.c

```
1    #include <stdio.h>
2    #include <string.h>
3    #include <unistd.h>
4    #include <uv.h>
5
6    uv_loop_t *loop;
7    uv_tty_t tty;
8    uv_timer_t tick;
9    uv_write_t write_req;
10   int width, height;
11   int pos = 0;
12   char *message = "  Hello TTY  ";
13
14   void update(uv_timer_t *req, int status) {
15       char data[500];
16
17       uv_buf_t buf;
18       buf.base = data;
19       buf.len = sprintf(data, "\033[2J\033[H\033[%dB\033[%dC\033[42;37m%s",
20                               pos,
21                               (width-strlen(message))/2,
22                               message);
23       uv_write(&write_req, (uv_stream_t*) &tty, &buf, 1, NULL);
24
25       pos++;
26       if (pos > height) {
```

```
27        uv_tty_reset_mode();
28        uv_timer_stop(&tick);
29    }
30 }
31
32 int main() {
33    loop = uv_default_loop();
34
35    uv_tty_init(loop, &tty, 1, 0);
36    uv_tty_set_mode(&tty, 0);
37
38    if (uv_tty_get_winsize(&tty, &width, &height)) {
39        fprintf(stderr, "Could not get TTY information\n");
40        uv_tty_reset_mode();
41        return 1;
42    }
43
44    fprintf(stderr, "Width %d, height %d\n", width, height);
45    uv_timer_init(loop, &tick);
46    uv_timer_start(&tick, update, 200, 200);
47    return uv_run(loop);
48 }
```

The escape codes are:

Code	Meaning
2 J	Clear part of the screen, 2 is entire screen
H	Moves cursor to certain position, default top-left
n B	Moves cursor down by n lines
n C	Moves cursor right by n columns
m	Obeys string of display settings, in this case green background (40+2), white text (30+7)

As you can see this is very useful to produce nicely formatted output, or even console based arcade games if that tickles your fancy. For fancier control you can try ncurses.

ABOUT

Nikhil Marathe started writing this book one afternoon (June 16, 2012) when he didn't feel like programming. He had recently been stung by the lack of good documentation on libuv while working on node-taglib. Although reference documentation was present, there were no comprehensive tutorials. This book is the output of that need and tries to be accurate. That said, Nikhil is young and inexperienced and may be severely wrong at points. He encourages you to call him out if you find an error. You can contribute to the book by forking the repository and sending a pull request.

Nikhil is indebted to Marc Lehmann's comprehensive man page about libev which describes much of the semantics of the two libraries.

The book is written in RestructuredText and uses Sphinx to generate the html and PDF. vim, tmux and iTerm2 on a Macbook Pro helped craft the text and code.

9.1 Licensing

The contents of this book are licensed as Creative Commons - Attribution. All code is in the **public domain**.

TEN

ALTERNATE FORMATS

The book is also available in:

- PDF

- epub